

Notes on the IP address-based redirection middleware

Srijan Technologies Pvt. Ltd., India
<http://www.srijan.in>

Document version: 1.0
July 6, 2009



Contents

| | | |
|----------|---|----------|
| 1 | Introduction | 2 |
| 2 | Installation and usage | 3 |
| 2.1 | External dependencies | 3 |
| 2.2 | Changes to <code>settings.py</code> | 3 |
| 2.3 | Running | 4 |
| 3 | Functioning of the middleware | 5 |
| 3.1 | Registered IP addresses | 5 |
| 3.2 | Unregistered IP addresses | 5 |
| 4 | Technical details | 6 |
| 5 | Issues | 8 |



Chapter 1

Introduction

This document describes the IP address redirection solution built for the Ciboe image database. The solution fulfils a need to have users from a set of registered IP addresses provided access to non-admin views without their having to log in. Administrative users from these IP addresses, namely anyone wishing to view a Django admin. page will need to log in as usual. For users from any unregistered IP addresses, the authentication process is as usual for Django, namely access is allowed if the user is authenticated and has sufficient permissions for the view in question.

The solution is implemented as Django middleware, also necessitating some changes to the URLs and views for the Django project. Chapter 2 describes installation and usage, and should be all that is required for a user of the middleware. Chapter 3 describes how the users of the system interact with it. Some technical details are provided in Chapter 4. Finally, policy decisions, edge cases, and peculiarities are described in Chapter 5.



Chapter 2

Installation and usage

2.1 External dependencies

The middleware uses the `netaddr` [3] library to manipulate IP addresses. `netaddr` was chosen after a careful examination of available Python IP libraries. The other strong contender was `ipaddr` [2] which has been accepted into the standard library for Python versions 2.7 and 3.1 onwards. The final choice of `netaddr` was because of better functionality, more features, and especially the end-user friendly classes like `IPRange` class for `IPAddress` ranges, and the `Wildcard` class. We use only a tiny part of the functionality provided by `netaddr`, and a conscious decision was made not to offer too many configuration options. However, if it is felt to be desirable, features like wildcards for IP addresses can be added. The current version of the library in use is 0.6.3.

2.2 Changes to `settings.py`

The following changes will need to be made to `settings.py` in order to use this middleware. All changes are actually made in `settings_local.py`:

- *Settings for middleware classes:* The `IPAddressMiddleware` resides in the file `middleware.py`, and can be added to the `MIDDLEWARE_CLASSES` variable in `settings.py` as `'middleware.IPAddressMiddleware'`. The `IPAddressMiddleWare` can be added any where in the list of middleware classes. Thus, the `MIDDLEWARE_CLASSES` variable could be set like:

```
MIDDLEWARE_CLASSES = (  
    'django.middleware.common.CommonMiddleware',  
    'django.contrib.sessions.middleware.SessionMiddleware',  
    'django.contrib.auth.middleware.AuthenticationMiddleware',  
    'middleware.IPAddressMiddleware',  
)
```

- *Specifying registered IPs:* This is given in the variable `SRJ_IP_REG` as a list, where each entry can be either a single IP address value specified as a string, or a sequence of two values as strings, which specify the start and end of an IP address range. Thus,



```
SRJ_IP_REG = [  
    '192.168.10.1', '192.168.10.4', ['192.168.10.6', '192.168.10.24'],  
    '192.168.10.40'  
]
```

Enhancements are also possible, e.g., the individual values can also themselves be ranges, e.g., '192.168.10.0/3'. For details, see the examples in <http://code.google.com/p/netaddr/wiki/IPv4Examples>. Still further extensions are possible, e.g., see <http://code.google.com/p/netaddr/wiki/WildcardExamples>, and the complete API at <http://packages.python.org/netaddr/>. However, in the interests of simplicity, these are not currently included. Anything that is not in the list of registered IP addresses is considered to be a non-registered IP address.

- *IPv6 vs. IPv4:* netaddr supports both IPv6 and IPv4 addresses, using similar syntax for both. So, IPv6 addresses should also work out of the box. While some static testing has been done, I do not have access to a IPV6 network in order to do a live test.
- *Templates for login pages:* There are separate login pages for users from registered, and unregistered IPs. These are configured through the variables SRJ_TMPL_LOGIN_REG and SRJ_TMPL_LOGIN_UNREG, e.g.,

```
SRJ_TMPL_LOGIN_REG = 'login_reg.html'  
SRJ_TMPL_LOGIN_UNREG = 'login.html'
```

Please note that these are set in the middleware, and hence specifying templates as `template_name` arguments to the dictionary for the login view will not work.

2.3 Running

After the above configuration, the only need should be to reload the webserver. For Apache2 on most Linux systems, this can be done with:

```
sudo /etc/init.d/apache2 reload
```



Chapter 3

Functioning of the middleware

This section describes the process followed by users accessing the site from the two classes of IP addresses: registered, and unregistered.

3.1 Registered IP addresses

Registered users are presented with login page A, which uses the SWF file, `maerskint.swf`. Normal users can just click on the Enter button, and get access to all non-admin views without the need to log in. Trying to access an admin. view, without having logged in earlier as an admin. user, will bring up a standard Django admin. login page. Administrative users can log in from the login page A, or can browse non-admin views till logging in through the Django admin. login page.

3.2 Unregistered IP addresses

Users from unregistered IP addresses function as usual with the Django authentication layer. They need to log in to access any pages protected with the `login_required` decorator, and need to have administrative privileges in order to access admin. views. A different login page, B, is shown to them, using the `maersk.swf` SWF file.



Chapter 4

Technical details

Below is a brief description of how the solution works. The middleware is in `middleware.py`, and uses only the `process_view` method:

- *Modification of views:* The views in `apps.files.views` have been modified, so that for every view function, say `view_func` which is protected by a `login_required` decorator, the decorator is removed, and “_reg” appended to the function name, e.g., `view_func_reg`. The normal function is then set up again with

```
view_func = login_required( view_func_reg )
```

Thus, the original functions are protected as before, and the new set of `XXX_reg` functions are not externally accessible as there are no URLs associated with them. The sole exception is the `login_reg` view which is used for logging in users from registered IP addresses. Access to this for users from non-registered IP addresses is prohibited in the middleware.

- *Login templates:* The middleware intercepts the login views for users from both registered and unregistered IP addresses, and provides the appropriate login template for each.
- *Denying access to login_reg for users from unregistered IPs:* The `login_reg` view is restricted to users from registered IP addresses. Requests for `login_reg` view from unregistered IP addresses are redirected to `login`.
- *Handling of non-admin views:* Non-admin. views from users coming from registered IP addresses are shown directly. A small complication here is that some views are from `apps.files.views` and others from `django.contrib.auth.admin.views`, so that the appropriate one needs to be found. Also, a small trick is that we directly call the view function from `viewfunc.__name__` whereas we are starting with an unregistered view. This works because we had set:

```
view_func = login_required( view_func_reg )
```

This assignment does not change `viewfunc.__name__`, so what we are really calling is the registered function, i.e., `view_func_reg` rather than `view_func`. This is a small matter, and we could have easily appended “_reg” if needed.



- *Handling of other views:* All other views are simply passed on, so that they are handled by the normal Django authentication layer. This includes admin. views from registered IP addresses, and all views from unregistered addresses.



Chapter 5

Issues

These issues have been noted earlier in email and Basecamp, and are provided here only for the sake of completeness.

- In `app/files/views.py`, for every `login_required` view, the `login_required` decorator has been removed, the view function name changed to `view_reg`, and a `login_required` version is retained with the original view name. The `view_reg` functions do not have configured URLs, and are hence not accessible. The only exception is `login_reg`, access to which is blocked from non-registered IP addresses by the middleware.
- Users after logging in were supposed to be redirected to `'/home/'`. This URL does not exist, so the redirection has been made to `''`
- Users from registered IPs looking at non-admin views cannot be authenticated (there is no user/password information). Hence the views should not have specific information shown for authenticated users.
- A user who has been authenticated from a registered (unregistered) IP, and switches to an unregistered (registered) IP is not required to authenticate again, and retains all earlier privileges.
- The login template is now set in the middleware, using the values in `settings.local.py`. This means that any values passed with `'template_name'` to the login view in `urls.py` will be overridden.
- Currently, in `urls.py`

```
(r'^admin/(. #)', include(admin.site.urls)),
```

is preferred for the admin. site configuration, instead of,

```
(r'^admin/(. #)', admin.site.root),
```

Please see <http://docs.djangoproject.com/en/dev/ref/contrib/admin/#hooking-adminsite-instances-into-your-urlconf>. With the old style declaration, I am forced to use

```
reverse(admin.site.root, args=[''])
```



to get the root of the admin. site URL. `admin.site.root` is deprecated and will disappear in Django 1.3. Now, the admin. uses named URLs, so that the usage should be

```
reverse('admin\_index', args=[''])
```

However, this will work only with the above changes to `urls.py`.

- `app.files.eps` is not protected by `login_required`. I am not sure why this is the case, and have left it alone.



Bibliography

- [1] A wiki write-up comparing netaddr to other Python IP address manipulation libraries. Though this is from the netaddr site, and might be biased, it does offer an argument for netaddr. <http://code.google.com/p/netaddr/wiki/YetAnotherPythonIPModule>.
- [2] ipaddr is an IP manipulation library in Python. <http://code.google.com/p/ipaddr-py/>.
- [3] netaddr is a Python library for the manipulation of various common network address notations and representations. <http://code.google.com/p/netaddr/>. netaddr is a super-set of the functionality in most of the Python IP address modules available on the net. See the Wiki note [1].

